

Bringing research to life

Innovation to facilitate data sharing in the life
sciences and biomedicine

NAOMI PENFOLD
Innovation Officer

Slides available at

<https://github.com/npscience/csvconfv3-presentation>



This work is licensed under a Creative Commons Attribution 4.0 International License.

About eLife

@eLifeInnovation

innovation@elifesciences.org



MAX-PLANCK-GESELLSCHAFT

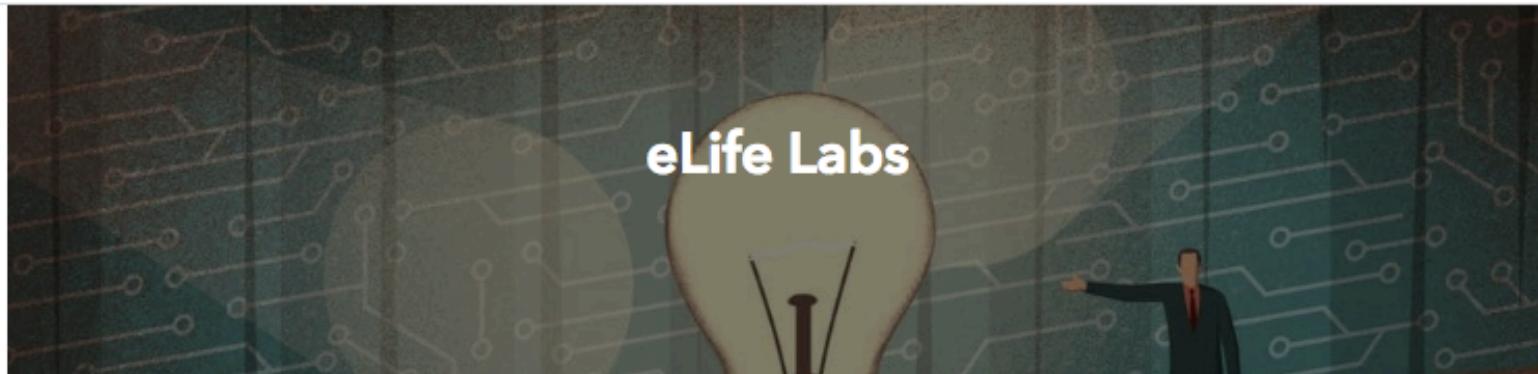
eLife is a non-profit organisation inspired by research funders and led by scientists

eLIFE

Helping scientists accelerate discovery by
operating a platform for research communication
that encourages and recognises the most
responsible behaviours in science

eLife Innovation Initiative

We invest in open source technologies, tools and processes that improve the way cutting-edge research is discovered, shared, consumed and evaluated



Exploring open-source solutions at the intersection of research and technology. Learn more about [innovation at eLife](#), or follow us on [Twitter](#).

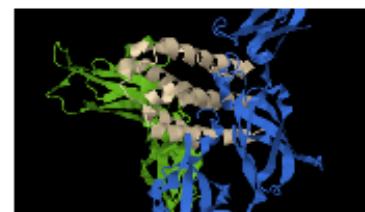
Latest



Composing reproducible manuscripts using R Markdown



Hack Cambridge Recuse entries: eXplore, Knowledge Direct, SciChat



Proteopedia for sharing macromolecule concepts online



The International Image Interoperability Framework (IIIF) for science publishers



The screenshot shows a web-based document editor with a dark-themed header. The header includes three dots on the left, the text "sciencefair" in a light font, and a close button with a black "X" on the right.

The main content area has a sidebar on the left labeled "Main Text". The main body contains two sections: "Background" and "Results".

The "Background" section contains the following text:

Many arthropod disease vectors have multiple opportunities to become infected with the same pathogen species during their lifetime (super-infection). The impact of super-infection within vectors to parasite transmission is largely unknown, and may have substantial impacts on epidemiology. For example, in the laboratory, pathogen transmission can be enhanced when different parasite species co-occur in the same individual vector, a phenomenon that has been observed in some [1 - 4] but not all mosquito species that have been tested [1, 4].

The "Results" section contains the following text:

The aim of this study was to investigate the potential epidemiological consequences of super-infection of mosquitoes by malaria parasites. Super-infection of vectors by successive parasite infections has been examined in a variety of infectious diseases [5 - 7], but to knowledge, the

On the right side, there is a vertical navigation bar with several sections:

- Contents
- Figures
- Info
- Background
- Methods
- Results
- Conclusions
- Main Text** (highlighted)
- Background
- Methods
- Mosquito feeds
- Statistical analysis
- Results
- Feeding behaviour following infection
- Re-exposure to parasites and infection
- Discussion
- Conclusions
- Authors' contributions

<https://github.com/codeforscience/sciencefair>

Ideas?

Email Naomi at innovation@elifesciences.org

Why share data+ ?

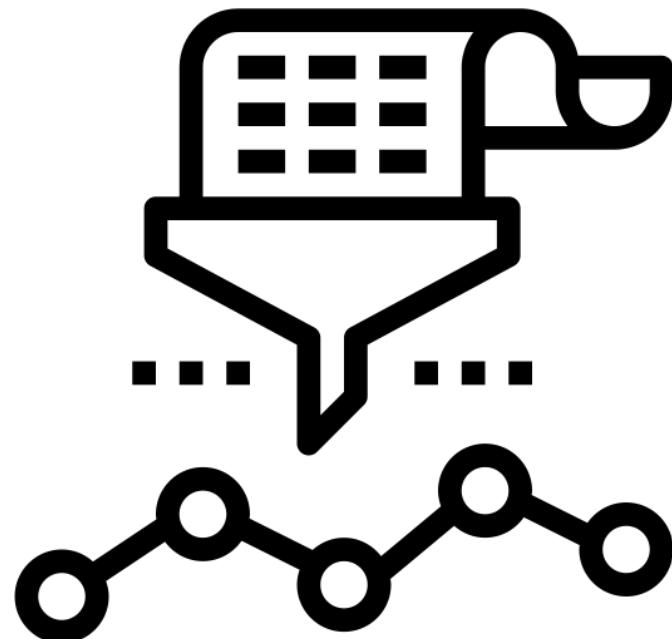
@eLifeInnovation

innovation@elifesciences.org

**“If I have seen further,
it is by **standing**
on the shoulders of giants”**

Isaac Newton

“Show me the evidence”

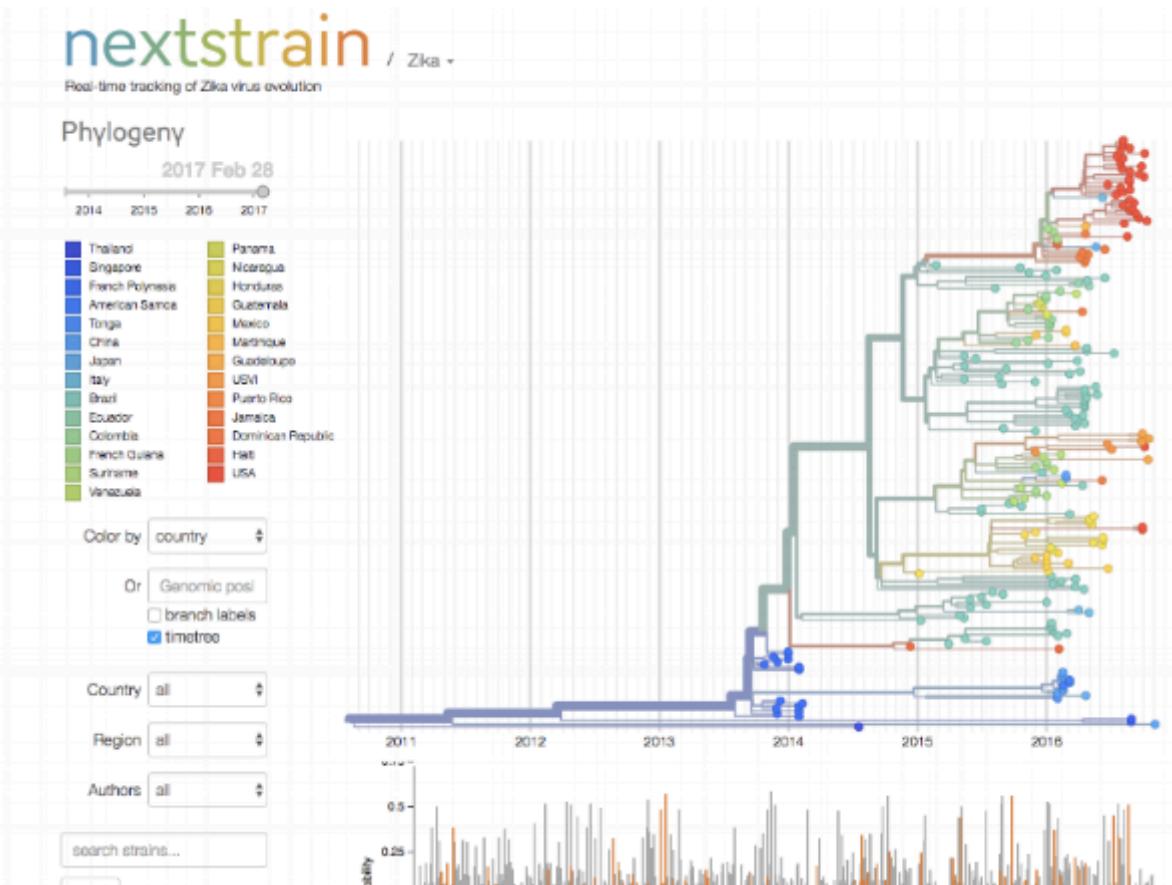


Created by Bebris
from Noun Project

**“True science thrives best
in glass houses, where
everyone can look in”**

Max Perutz

Sharing data accelerates discovery and innovation



Sharing data allows community to spot problems...

Gene name errors are widespread in the scientific literature

Mark Ziemann, Yotam Eren and Assam El-Osta 

Genome Biology 2016 17:177 | DOI: 10.1186/s13059-016-1044-7 | © The Author(s). 2016

Published: 23 August 2016

- 'SEPT2' (Septin 2) → '2-Sep'
- '2310009E13' (RIKEN identifier) → '2.31E+13'

1 in 5

...and adjust with advancements

Brains

WIRED

Bug in fMRI software calls 15 years of research into question

Popular pieces of software for fMRI were found to have false positive rates up to 70%

By EMILY REYNOLDS

Wednesday 6 July 2016

The state of Open Data

@eLifeInnovation

innovation@elifesciences.org

Open... and FAIR?

F Findable

A Accessible

I Interoperable

R Re-usable

<https://www.force11.org/group/fairgroup/fairprinciples>

Recent surveys



Van den Eynden, Veerle et al. (2016)
Towards Open Research: Practices,
experiences, barriers and
opportunities. Wellcome Trust. <https://dx.doi.org/10.6084/m9.figshare.4055448>

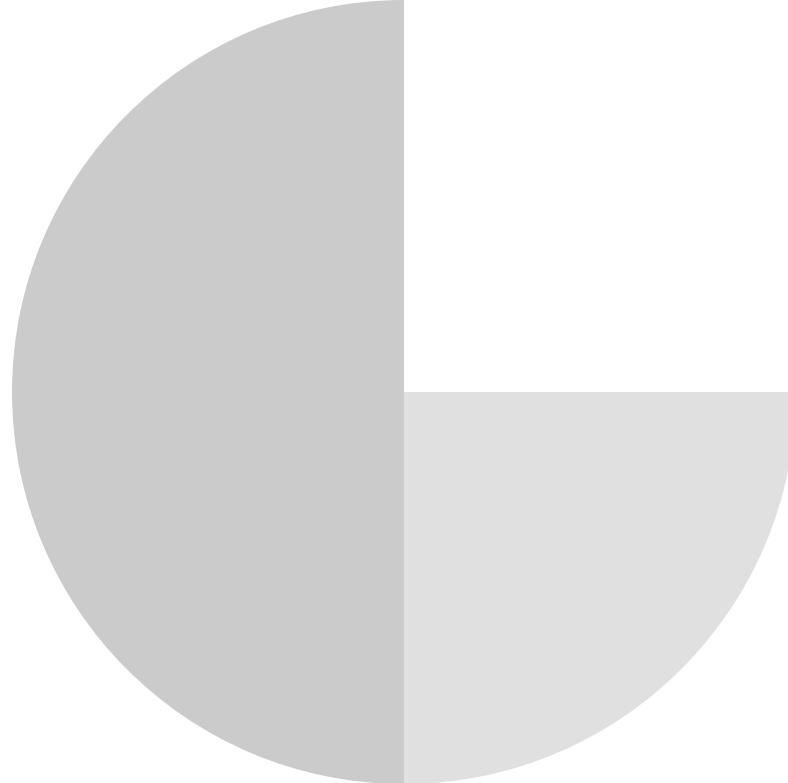


Treadway, Jon; Hahnel, Mark; Leonelli, Sabina; Penny, Dan; Groenewegen, David; Miyairi, Nobuko; Hayashi, Kazuhiro; O'Donnell, Daniel; Science, Digital; Hook, Daniel (2016): The State of Open Data Report. figshare. <https://doi.org/10.6084/m9.figshare.4036398.v1>

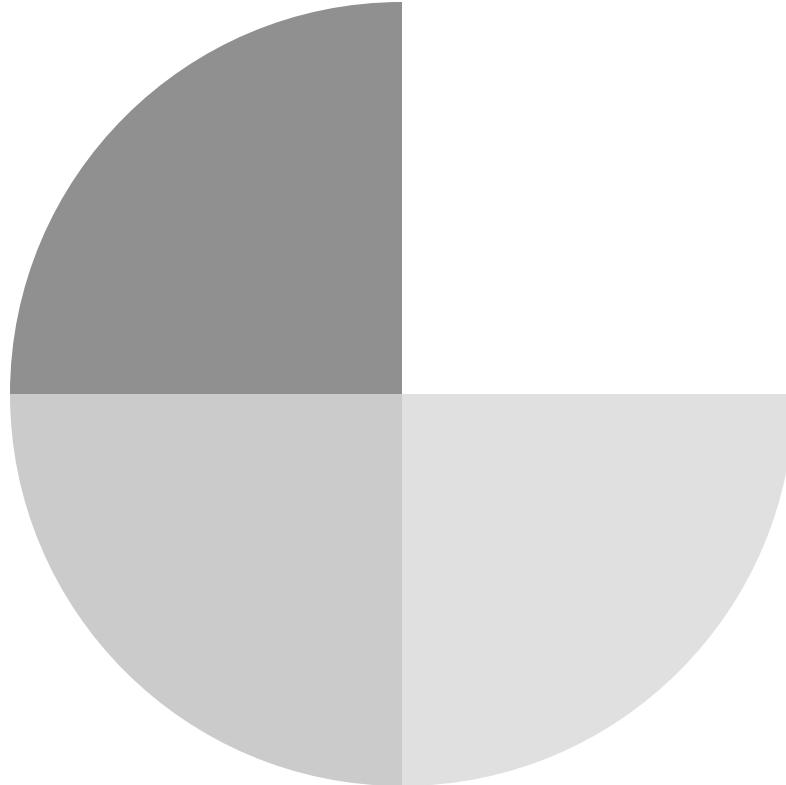
73% of biologists aware



50% make data available



25% make full dataset available



Improving data+ sharing

@eLifeInnovation

innovation@elifesciences.org

Data sharing policies

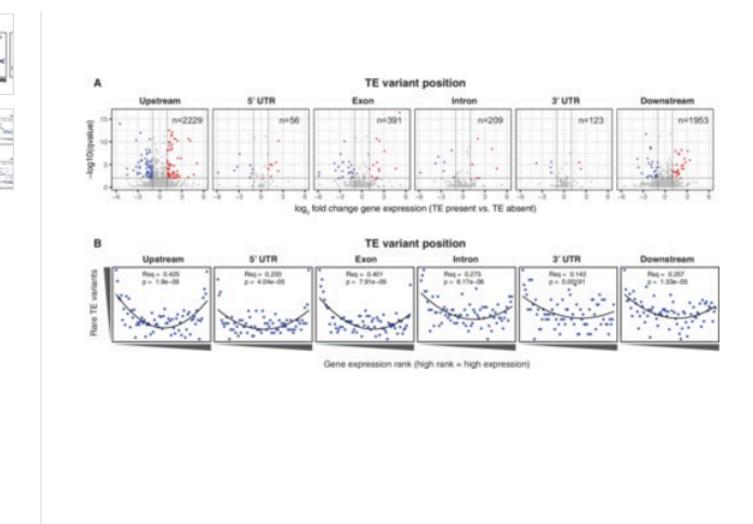


Figure 4.

[Download figure](#)

[Open in new tab](#)

Differential transcript abundance associated with TE variant presence/absence.

(A) Transcript abundance differences for genes associated with TE insertion variants at different positions, indicated in the plot titles. Genes with significantly different transcript abundance in accessions with a TE insertion compared to accessions without a TE insertion are colored blue (lower transcript abundance in accessions containing TE insertion) or red (higher transcript abundance in accessions containing TE insertion). Vertical lines indicate ± 2 fold change in FPKM. Horizontal line indicates the 1% false discovery rate. (B) Relationship between rare TE variant counts and gene expression rank. Cumulative number of rare TE variants in equal-sized bins for gene expression ranks, from the lowest-ranked accession (left) to the highest-ranked accession (right). Lines indicate the fit of a quadratic model.

DOI: <http://dx.doi.org/10.7554/eLife.20777.025>

Figure 4—source data 1.

Differentially expressed genes associated with TE presence/absence.

List of genes differentially expressed dependent on the presence/absence of nearby TE variants.

DOI: <http://dx.doi.org/10.7554/eLife.20777.026>

[Download source data \[figure-4—source-data-1.media-6.tsv\]](#)

Publish

Describe

Interact

Reproduce

Update

Link

Recognise

Data publications

(GIGA)ⁿ
SCIENCE

SCIENTIFIC DATA

110110
0111101
11011110
011101101

Publish

Describe

Interact

Reproduce

Update

Link

Recognise

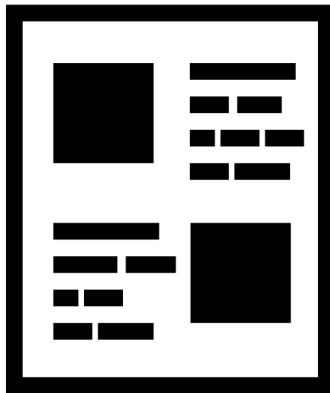


elifesciences.org

@eLifeInnovation

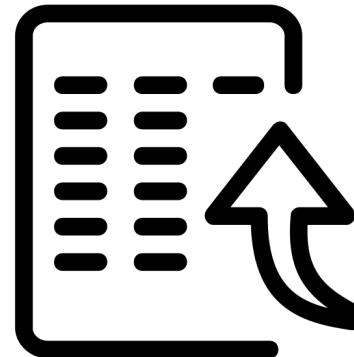
25

Can we bring the data closer to the narrative?



Created by Dmitry Podluzny
from Noun Project

+



Created by IconfactoryTeam
from Noun Project

The interactive figure



Au Data-driven, interactive scienc eLife

Secure https://www.authorea.com/users/3/articles/3904-data-driven-interactive-science-with-d3.js-plots-and-ipython-notebooks/_s...

[ABOUT](#) [EXPLORE](#) [HELP](#) [LOG IN](#) [SIGN UP](#)

[☆](#) [f](#) [v](#) [p](#)

Data-driven, interactive science, with d3.js plots and IPython Notebooks

 Alberto Pepe
 Nathan Jenkins
 Matteo Cantiello

1 Javascript, d3.js, and d3po.js

4

Javascript offers many ways to create data-driven graphics. A popular solution is [D3.js](#), a JavaScript library to create and control web-based dynamic and interactive graphical forms. A gallery of some beautiful d3.js plots can be found [here](#).

Authorea now supports most Javascript-based data visualization solutions. The example below - Figure 1 - is a plot generated using [D3po.js](#) which is a javascript extension of d3.js. D3po allows anyone with no special data visualization skills, to make an interactive, publication-quality figure that has staged builds and linked brushing through scatter plots. What's even cooler is that the plot below is based on actual data (astrophysics data, yay!). The figure describes how metallicity affects color in cool stars. It is based on work of graduate student Elizabeth Newton and others ([Newton 2014](#)). Try clicking and dragging in the scatter plots to understand the power of linked brushing in published figures.

You should know that this entire visualization is running within Authorea. The Javascript, HTML, CSS and all the data associated with this image are all part of this blog post. They are individual files which can be found by clicking on the folder icon on the top left corner of this page.

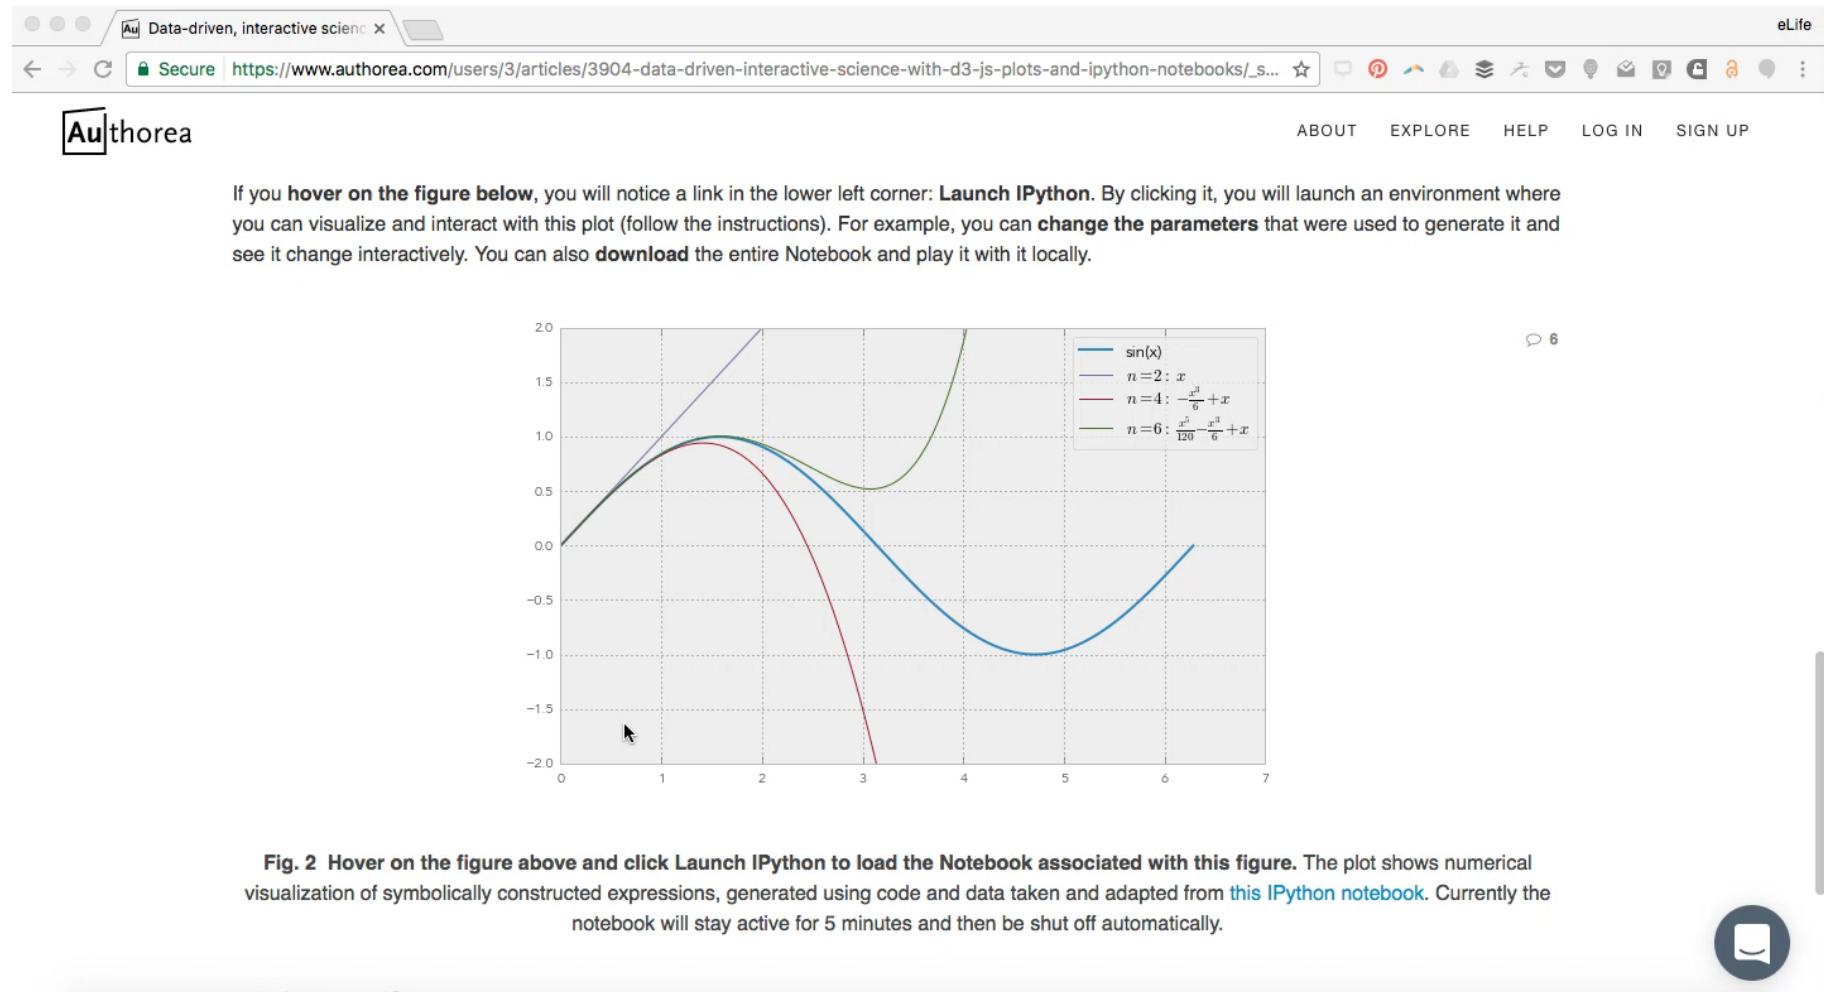
1

Some steps removed and sequences shortened

The executable figure

Some steps removed and sequences shortened





Some steps removed and sequences shortened

A screenshot of a web browser window titled "binder". The address bar shows "mybinder.org". The page content features the Binder logo (three overlapping circles in orange, blue, and red) and the text: "Turn a GitHub repo into a collection of interactive notebooks powered by Jupyter and Kubernetes." Below this, a paragraph explains that users can add a badge to their GitHub repository to make their code reproducible. A call-to-action button says "Tell us your GitHub repo" with a placeholder "user/project OR github url". To the right, a note specifies that the URL should point to a Jupyter notebook named "index.ipynb".

binder

mybinder.org

Jeremy

binder

Turn a GitHub repo into a collection of interactive notebooks powered by Jupyter and Kubernetes.

Have a repo full of Jupyter notebooks? With Binder, you can add a badge that opens those notebooks in an executable environment, making your code immediately reproducible by anyone, anywhere.

100% free and open source. Browse examples. Read the FAQ.
(Currently in testing, let us know if you run into trouble!)

1

Tell us your GitHub repo

user/project OR github url

This should contain Jupyter notebooks. If one of them is called index.ipynb it will be where your Binder starts. Any extra folders or files (e.g. data) will be included. See an [example](#) repo that uses Binder.

The reproducible document

The screenshot shows the RStudio interface. On the left, the code editor displays an R Markdown file named 'test.rmd' with the following content:

```
1 ---  
2 title: "Blog example"  
3 author: "CHJ Hartgerink"  
4 date: "March 30, 2017"  
5 output: html_document  
6 ---  
7  
8 # R Markdown  
9  
10 This is a test result, t(69) = 1.95, p = `r  
round(pt(q = 1.95, df = 69, lower.tail = FALSE),  
3)`|  
11
```

On the right, the rendered output is shown in the 'Viewer' pane:

Blog example
CHJ Hartgerink
March 30, 2017

R Markdown

This is a test result, t(69) = 1.95, p = 0.028



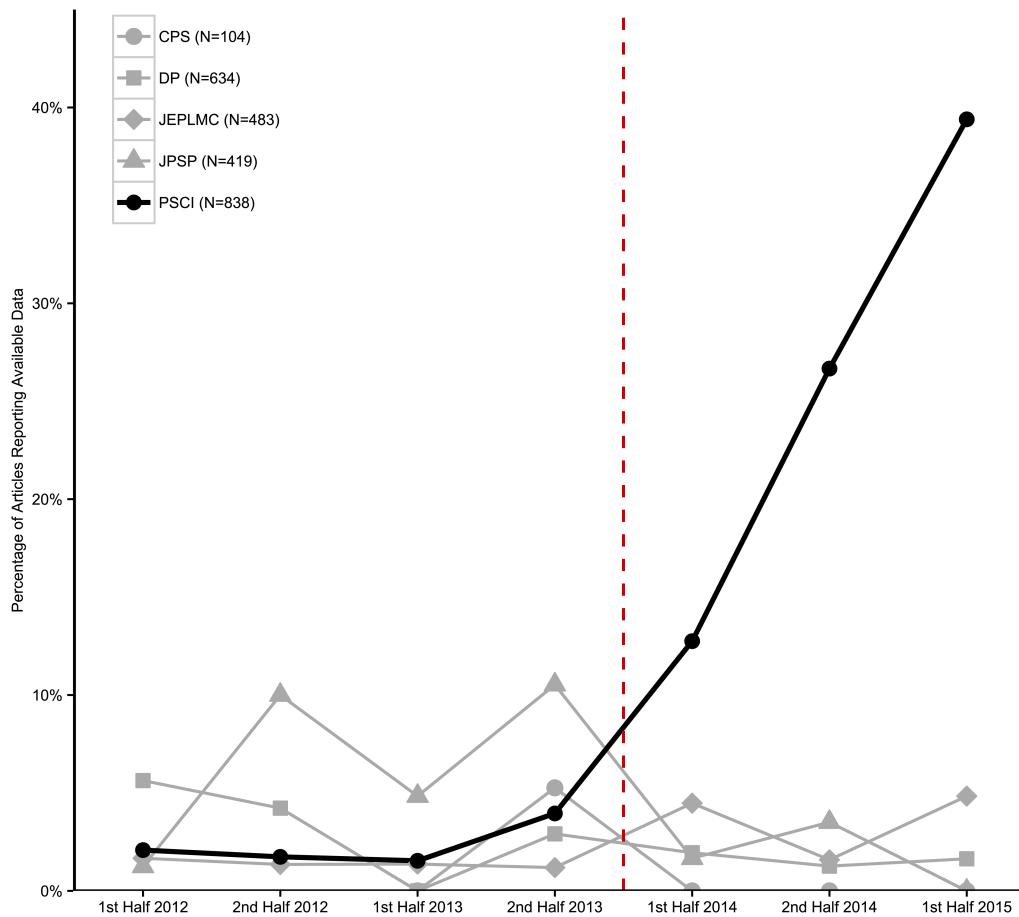
"It took me a couple of hours to...
REPRODUCE EXACTLY the analysis presented in the manuscript...
With few more hours, I **was able to modify the authors' code**
to change a linear scale for a log scale for their Fig. 4."

Christophe Pouzat, reviewer
GigaScience blog: <http://gigasciencejournal.com/blog/qa-on-dynamic-documents>

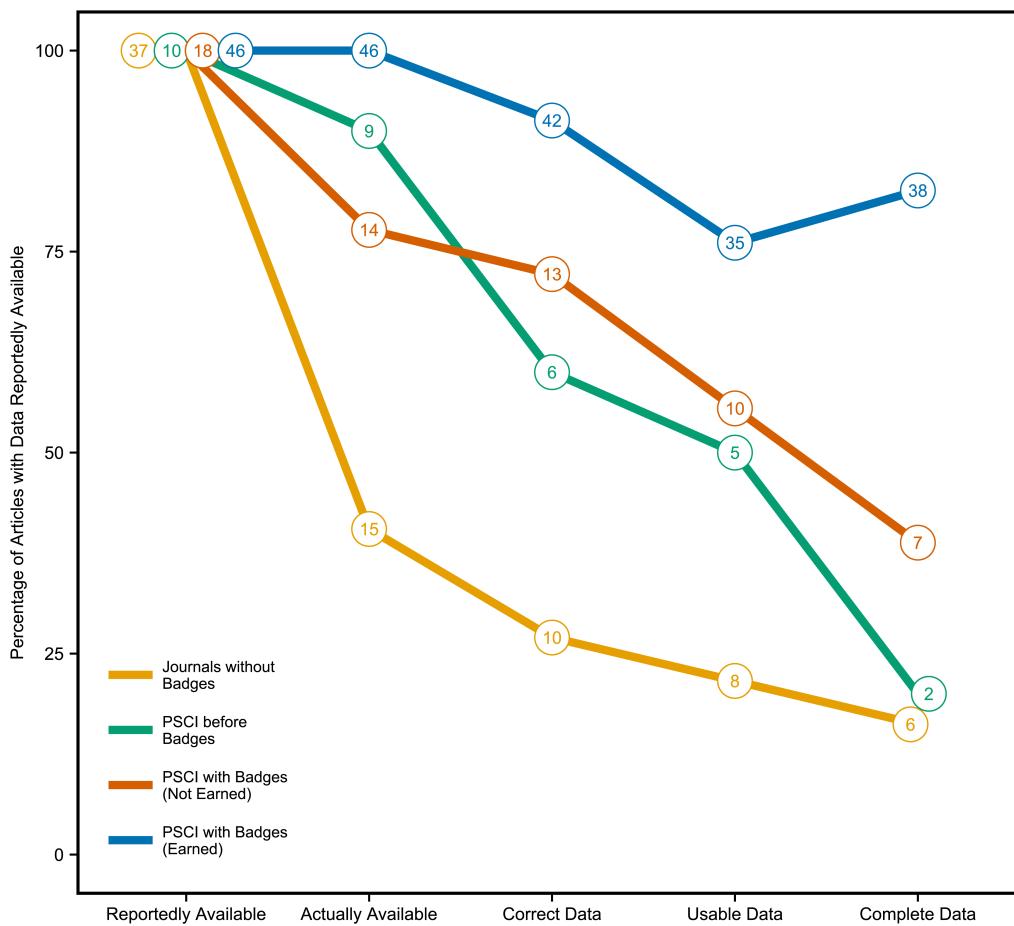


Credit





PLOS Biology, 2016; <https://doi.org/10.1371/journal.pbio.1002456.g002>



PLOS Biology, 2016; <https://doi.org/10.1371/journal.pbio.1002456.g002>

The screenshot shows a web browser window with the title bar "LCAV Reproducible Research". The address bar contains the URL "rr.epfl.ch/paper/MUT2016". The page header includes the LCAV logo, navigation links for "Home", "Browse", "Resources", and "Stats", and a "Sign in" button. The main content area displays a research paper summary.

Where You Are Is Who You Are: User Identification by Matching Statistics

Farid Movahedi Naini & Jayakrishnan Unnikrishnan & Patrick Thiran & Martin Vetterli

IEEE Journal of Selected Topics in Signal Processing

Most users of online services have unique behavioral or usage patterns. These behavioral patterns can be used to identify and track users by using only the observed patterns in the behavior. We study the task of identifying users from statistics of their behavioral patterns. Specifically, we focus on the setting in which we are given histograms of users' data collected in two different experiments. In the first dataset, we assume that the users' identities are anonymized or hidden and in the second dataset we assume that their identities are known. We study the task of identifying the users in the first dataset by matching the histograms of their data with the histograms from the second dataset. In a recent work [1], [2] the optimal algorithm for this user identification task was introduced. In this paper, we evaluate the effectiveness of this method on a wide range of datasets with up to 50, 000 users, and in a wide range of scenarios. Using datasets such as call data records, web browsing histories, and GPS trajectories, we demonstrate that a large fraction of users can be easily identified given only histograms of their data, and hence these histograms can act as users' fingerprints. We also show that simultaneous identification of users achieves better performance compared to one-by-one user identification. Furthermore, we show that using the optimal method for identification does indeed give higher identification accuracies than heuristics-based approaches in such practical scenarios. The accuracies obtained under this optimal method can thus be used to quantify the maximum level of user identification that is possible in such settings. We show that the key factors affecting the accuracy of the optimal identification algorithm are the duration of the data collection, the number of users in the anonymized dataset, and the resolution of the dataset. We also analyze the effectiveness of k-anonymization in resisting user identification attacks on these datasets.

2016

[Download PDF](#)



Code

[Rate the code](#)

Easy to reproduce figures



[Download Data](#)

[Cite this paper](#)



Publish

Describe

Interact

Reproduce

Update

Recognise

Link

Where You Are Is Who You Are: User Identification by Matching Statistics

Farid Movahedi Naini & Jayakrishnan Unnikrishnan & Patrick Thiran & Martin Vetterli

IEEE Journal of Selected Topics in Signal Processing

Most users of online services have unique behavioral or usage patterns. These behavioral patterns can be used to identify and track users by using only the observed patterns in the behavior. We study the task of identifying users from statistics of their behavioral patterns. Specifically, we focus on the setting in which we are given histograms of users' data collected in two different experiments. In the first dataset, we assume that the users' identities are anonymized or hidden and in the second dataset we assume that their identities are known. We study the task of identifying the users in the first dataset by matching the histograms of their data with the histograms from the second dataset. In a recent work [1], [2] the optimal algorithm for this user identification task was introduced. In this paper, we evaluate the effectiveness of this method on a wide range of datasets with up to 50, 000 users, and in a wide range of scenarios. Using datasets such as call data records, web browsing histories, and GPS trajectories, we demonstrate that a large fraction of users can be easily identified given only histograms of their data, and hence these histograms can act as users' fingerprints. We also show that simultaneous identification of users achieves better performance compared to one-by-one user identification. Furthermore, we show that using the optimal method for identification does indeed give higher identification accuracies than heuristics-based approaches in such practical scenarios. The accuracies obtained under this optimal method can thus be used to quantify the maximum level of user identification that is possible in such settings. We show that the key factors affecting the accuracy of the optimal identification algorithm are the duration of the data collection, the number of users in the anonymized dataset, and the resolution of the dataset. We also analyze the effectiveness of k-anonymization in resisting user identification attacks on these datasets.

2016

[Download PDF](#)



Code

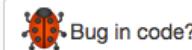
[Rate the code](#)

Easy to reproduce figures



[Download Data](#)

[Cite this paper](#)



Bug in code?

Please share your thoughts with us

How easy it was to reproduce figures?



Was the code well-commented?



Was the code easy to apply to new scenarios?



Your comments are really appreciated

Your email (optional)

Send

I'm not a robot



Privacy - Terms

Close

Submit

Where You Are Is Visible from Matching Statistics

Farid Movahedi Naini & Jayakrishnan Venkateswaran

IEEE Journal of Selected Topics in Signal Processing

Most users of online services have unique identifiers that are used to identify and track users by using identifying users from statistics of their data. Even if user identities are given histograms of users' data collected over time, the users' identities are anonymized or completely unknown. We study the task of identifying users from histograms with the histograms from the second dimension. This identification task was introduced. In this work, we consider datasets with up to 50,000 users, and show that users can be easily identified given only histograms of their device fingerprints. We also show that simultaneous one-by-one user identification. Furthermore, we show that histograms indeed give higher identification accuracy. The accuracies obtained under this optimal identification that is possible in such scenarios. The accuracies of the optimal identification algorithm are the dataset size, the resolution of the dataset, and the resolution of the data. User identification attacks on these datasets are also discussed.

2016

Download PDF



Code

Rate the code

Easy to reproduce figures



Download Data

Cite this paper



Bug in code?

easily identified given only histograms of their data, and hence these histograms can act as users' fingerprints. We also show that simultaneous identification of users achieves better performance compared to one-by-one user identification. Furthermore, we show that using the optimal method for identification does indeed give higher identification accuracies than heuristics-based approaches in such practical scenarios. The accuracies obtained under this optimal method can thus be used to quantify the maximum level of user identification that is possible in such settings. We show that the key factors affecting the accuracy of the optimal identification algorithm are the duration of the data collection, the number of users in the anonymized dataset, and the resolution of the dataset. We also analyze the effectiveness of k-anonymization in resisting user identification attacks on these datasets.

[Download Data](#)[Cite this paper](#)

Supplementary Materials

How to run the code

First, download the code and data files, uncompress them and copy/paste the content of the data files into the `generatingResultsFigures/` folder. Then, to generate the figures in the paper, simply open **MATLAB** and navigate to the `generatingResultsFigures/` folder (in the code files) and follow the instructions below.

Note that since only the *WBH* and the *GL* datasets are publicly available (in [here](#) and [here](#), respectively), the codes generate only the figures related to these two datasets.

Figure 8 and 9

In MATLAB, type

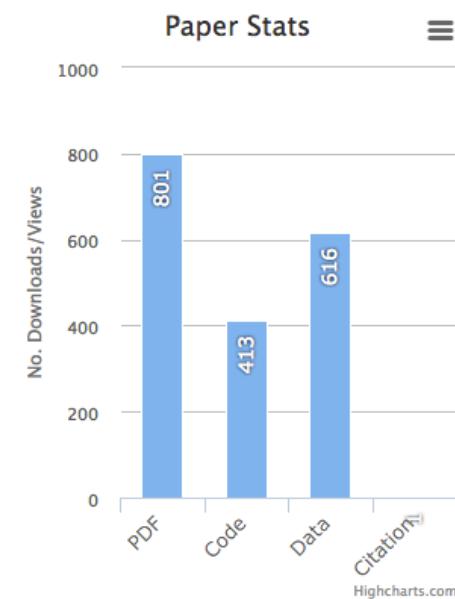
```
>> WBH_results1.m
```

and

```
>> WBH_results2.m
```

to reproduce Figures 8 and 9, respectively.

Figure 10 and 11



Connecting projects

- Linked open data
- Human curation



ReFigure



Figures hyperlinked to original paper

<basic navigation menu>

Report/Flag

Longest post onset detection of Zika RNA in semen in multiple case studies

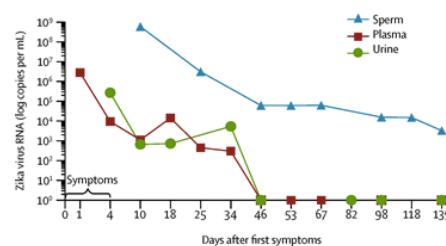
Christa Osuna

| Type of test and sample | Results | | | |
|------------------------------|-------------------|-------------------|-------------------|-------------------|
| | Day 12* | Day 93* | Day 134* | Day 188* |
| ZIKV real-time RT-PCR serum | Neg | Neg | Neg | NT |
| ZIKV real-time RT-PCR urine | Neg | Pos (Ct: 36.3) | Neg | NT |
| ZIKV real-time RT-PCR saliva | Neg (Ct: 36.4) | Pos (Ct: 36.4) | Neg | NT |
| ZIKV real-time RT-PCR semen | NT | Pos (Ct: 29.6) | Pos (Ct: 32.5) | Pos (Ct: 30.2) |
| IFA ZIKV IgM titre | 1:60 | 1:40 | 1:20 | 1:20 |
| IFA ZIKV IgG titre | 1:60 | 1:20 | 1:20 | 1:40 |
| MNT antibody titre | 1:60 | at 320 | at 320 | NT |

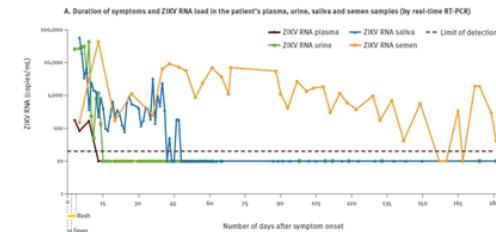
*Number of days after symptom onset.

Laboratory findings related to Zika virus infection in a traveller returning from Haiti to Italy, Feb-July 2016

Nicostri et al
Eurosurveillance



Zika virus in semen and spermatozoa
Mansuy et al
Lancet Infectious Diseases



Clinical and laboratory findings in a patient returning from Haiti to Italy, Jan 2016
Barzon et al
Euro Surveillance



<http://refigure.org/index.html>

Where next?

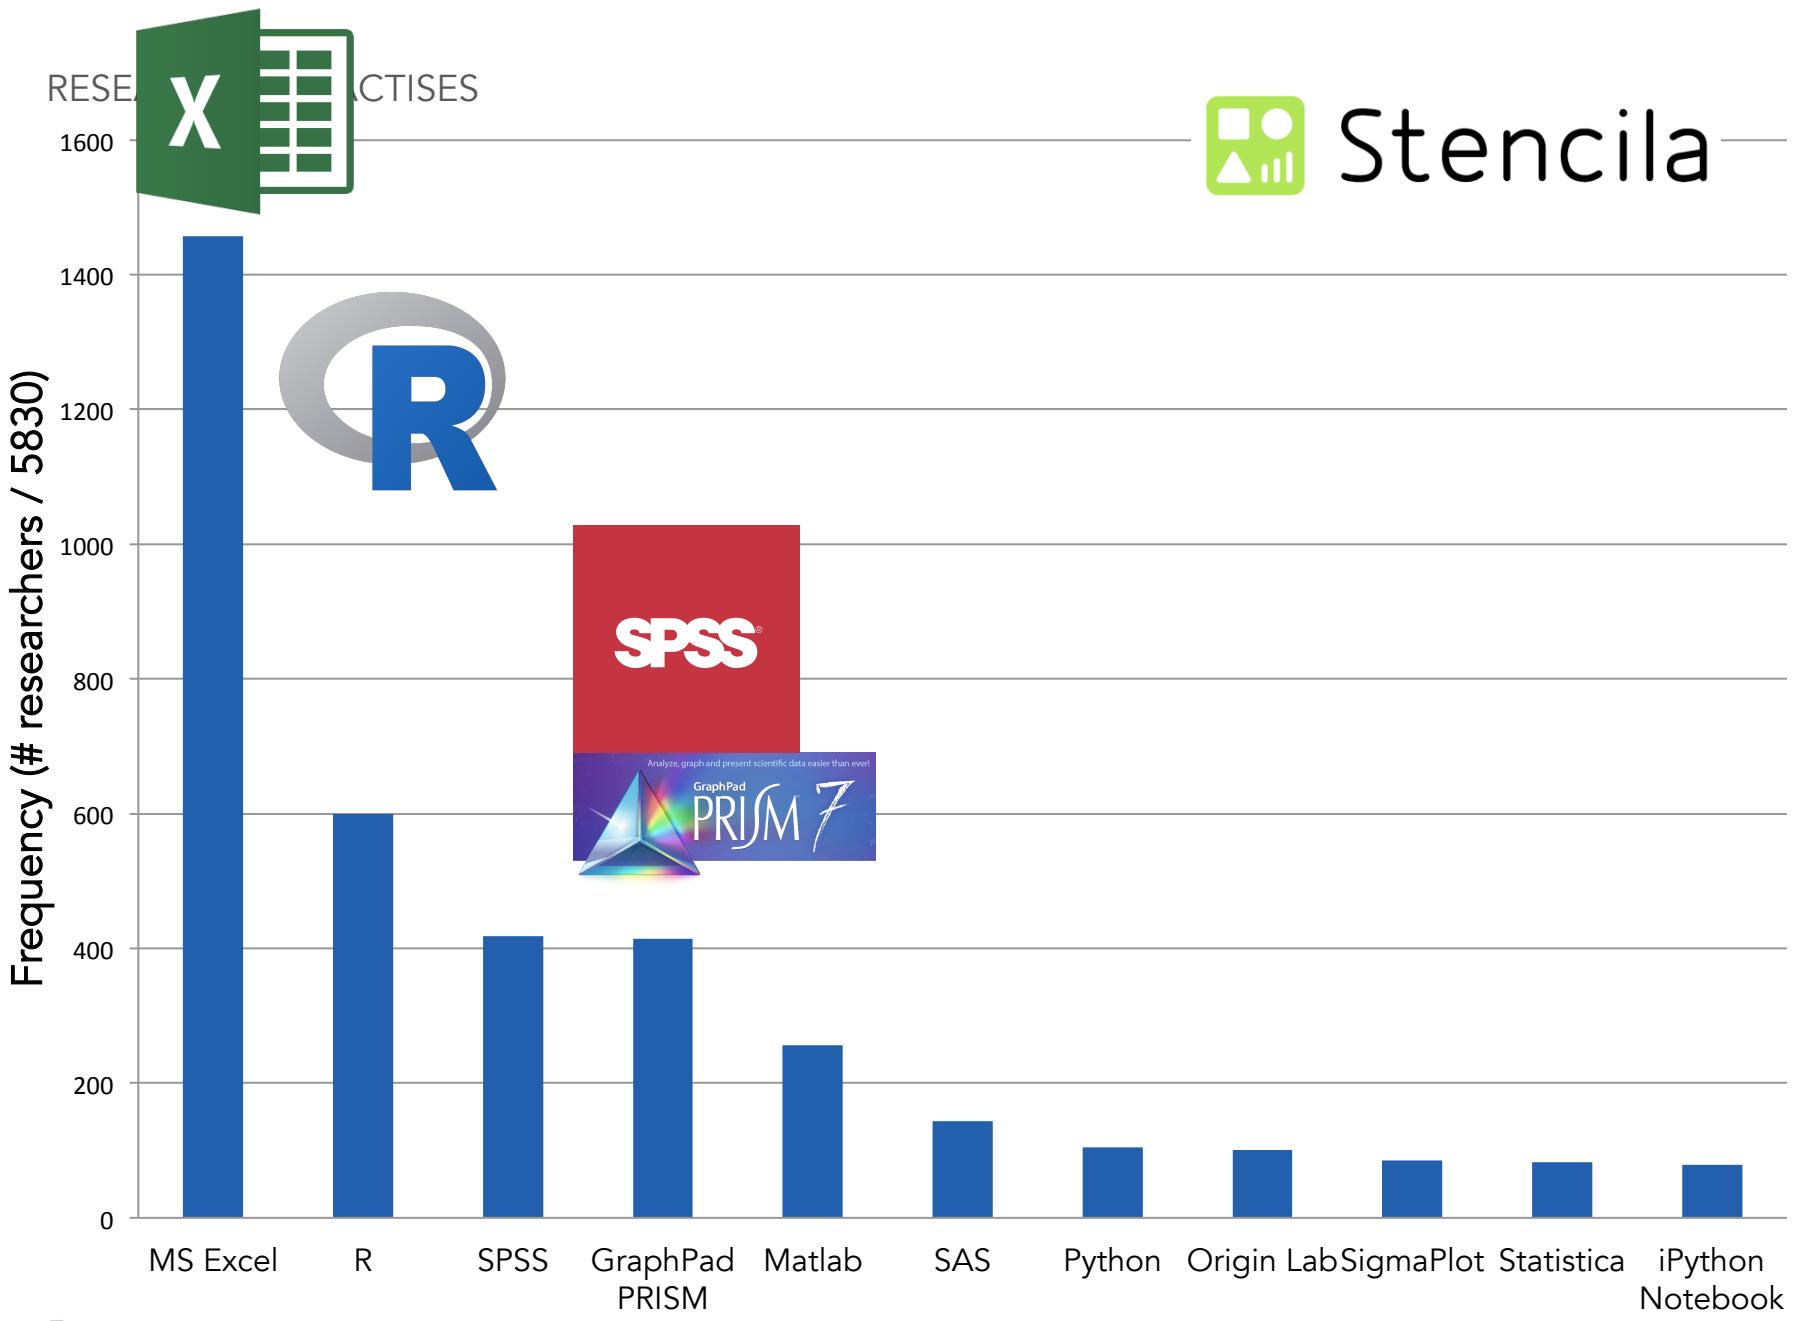
@eLifeInnovation

innovation@elifesciences.org

Publisher constraints

- Static version of peer-reviewed article
- Persistence
- Accuracy
- Who stores? Who hosts? Who computes?

...and researcher practises



Outstanding issues

- Improve quality of data sharing
- Incentives

“No time or funding” to learn new tools, document data well

“No benefit”

Opportunities

- Facilitating data sharing **via the journal** means:
 - Persistence
 - Leverages incentive system
 - Streamline the process from initial discovery to replication
 - Engages community
- Doing it **open source** means:
 - Progress can be shared
 - Stimulates further innovation
 - Engages community

Your thoughts?

Naomi Penfold, Innovation Officer

@eLifeInnovation

Email: innovation@elifesciences.org

Slides: github.com/npscience/csvconfv3-presentation

“The **impact we cherish is
discovery in science”**

Randy Schekman, eLife Editor-in-Chief